

# IERG5050 AI Foundation Models, Systems & Applications

## Fall 2025

### Agentic AI Applications

Prof. Wing C. Lau

[wclau@ie.cuhk.edu.hk](mailto:wclau@ie.cuhk.edu.hk)


<http://www.ie.cuhk.edu.hk/wclau>

# Acknowledgements


Many of the slides in this lecture are adapted from the sources below. Copyrights belong to the original authors.

- LangChain for LLM Application Development – a short course by Harrison Chase, Andrew Ng, <https://www.deeplearning.ai/short-courses/langchain-for-llm-application-development/>
- LangChain Chat with your Data – a short course by Harrison Chase, <https://www.deeplearning.ai/short-courses/langchain-chat-with-your-data/>
- AI Agents in LangGraph– a short course by Harrison Chase, Rotem Weiss, <https://www.deeplearning.ai/short-courses/ai-agents-in-langgraph/>
- Building LLM Powered Applications By Valentina Alto, Packt Publishing, May 2024
- AI Agents in Action, by Michael Lanham, Feb 2025, Manning Publications
- LLMs in Production by Christopher Brousseau, Matthew Sharp, Jan 2025, Manning Publications
- Shunyu Yao, “LLM agents: brief history and overview,” Guest Lecture for UC Berkeley MOOC on Large Language Model Agents, Fall 2024, [https://rdi.berkeley.edu/llm-agents/assets/llm\\_agent\\_history.pdf](https://rdi.berkeley.edu/llm-agents/assets/llm_agent_history.pdf), <https://www.youtube.com/watch?v=RM6ZArd2nVc>
- Chi Wang, “Agentic AI Frameworks & AutoGen,” Guest Lecture for UC Berkeley MOOC on Large Language Model Agents, Fall 2024, <https://rdi.berkeley.edu/llm-agents-mooc/slides/autogen.pdf>, <https://www.youtube.com/watch?v=OQdlmCMSOo4>
- Jerry Liu, “Building a Multimodal Knowledge Assistant (LlamaIndex),” Guest Lecture for UC Berkeley MOOC on Large Language Model Agents, Fall 2024, <https://rdi.berkeley.edu/llm-agents-mooc/slides/MKA.pdf>, <https://www.youtube.com/watch?v=OQdlmCMSOo4>
- Ruslan Salakhutdinov, “Multimodal Autonomous AI Agents,” Guest Lecture for UC Berkeley MOOC on Advanced Large Language Model Agents, Spring 2025, <https://rdi.berkeley.edu/adv-llm-agents/slides/ruslan-multimodal.pdf>, <https://www.youtube.com/live/RPINOYM12RU>
- Nicolas Chapados, “AI Agents for Enterprise Workflows,” Guest Lecture for UC Berkeley MOOC on Large Language Model Agents, Fall 2024, <https://rdi.berkeley.edu/llm-agents/assets/agentworkflows.pdf>, <https://www.youtube.com/live/-yf-e-9FvOo>
- Ben Mann, “Measuring Agent capabilities and Anthropic’s RSP,” Guest Lecture for UC Berkeley MOOC on Large Language Model Agents, Fall 2024, <https://rdi.berkeley.edu/llm-agents/assets/antrsp.pdf>, <https://www.youtube.com/live/6v2AnWolZoo>
- Caiming Xiong, “Multimodal Agents – from Perception to Action,” Guest Lecture for UC Berkeley MOOC on Advanced Large Language Model Agents, Spring 2025, [https://rdi.berkeley.edu/llm-agents-mooc/slides/Multimodal\\_Agent\\_caiming.pdf](https://rdi.berkeley.edu/llm-agents-mooc/slides/Multimodal_Agent_caiming.pdf), [https://www.youtube.com/live/n\\_\\_Tim8K2IY](https://www.youtube.com/live/n__Tim8K2IY)
- Omar Khattab, “Compound AI Systems & the DSPy Framework,” Guest Lecture for UC Berkeley MOOC on Large Language Model Agents, Fall 2024, <https://rdi.berkeley.edu/llm-agents-mooc/slides/dspy Lec.pdf>, <https://www.youtube.com/live/JEMYuzrKLUw>
- Graham Neubig, “Agents for Software Development,” Guest Lecture for UC Berkeley MOOC on Large Language Model Agents, Fall 2024, <https://rdi.berkeley.edu/llm-agents-mooc/slides/neubig24softwareagents.pdf>, <https://www.youtube.com/live/f9L9Fkg-8Kd>

# An LLM-powered Application




# Process-Flow of a Sample LLM-powered Application: a Minimalist ChatBot







# Copilot (aka AI Assistant): An LLM-powered Application




An LLM-powered Copilot




A Copilot w/ grounding



A conversational UI to reduce the gap between the user and the database or ext. tools




# A Multimodal Application via an LLM w/ Single-modal Tools



*To contrast with a Multimodal-NATIVE LLM (later)*


# High-level Architecture of an LLM-powered Application



# LLM-application/ Agent Development Frameworks


| Framework                 | Best For                                  | Deployment                            | Scalability                        | Current Adoption              |
|---------------------------|-------------------------------------------|---------------------------------------|------------------------------------|-------------------------------|
| LangChain & LangGraph     | Enterprise AI apps, multi-agent workflows | Self-hosted, Cloud                    | High – widely used in production   | Most adopted                  |
| Autogen & Semantic Kernel | Microsoft ecosystem, scalable AI apps     | Azure, Self-hosted                    | High – enterprise-scale            | Growing in enterprises        |
| LlamaIndex                | AI-powered search, RAG                    | Cloud, Self-hosted                    | High – efficient data processing   | Strong in AI search           |
| AutoGPT                   | No-code AI automation, continuous agents  | Cloud-first, some self-hosted options | Medium – automation focused        | Popular for prototyping       |
| CrewAI                    | Workflow-based multi-agent AI apps        | Cloud, Self-hosted                    | Medium – still early-stage         | Fast-growing                  |
| PydanticAI                | FastAPI & Pydantic-based AI apps          | Self-hosted                           | Low – lightweight framework        | Niche adoption                |
| Spring AI                 | Java-based AI applications                | Self-hosted, Cloud                    | Medium – depends on Java ecosystem | Popular in Java world         |
| Haystack                  | LLM-powered search & RAG                  | Self-hosted, Cloud                    | High – built for production        | Strong in search applications |

# LangChain: a framework for building LLM-based Applications




Components of LangChain

# Haystack: a framework for building LLM-based Applications



Haystack's components

# Semantic Kernel: a framework for building LLM-based Applications



## Anatomy of Semantic Kernel

# LangChain vs. Haystack vs. Semantic Kernel

| Feature               | LangChain                   | Haystack                    | Semantic Kernel             |
|-----------------------|-----------------------------|-----------------------------|-----------------------------|
| LLM support           | Proprietary and open-source | Proprietary and open source | Proprietary and open source |
| Supported languages   | Python and JS/TS            | Python                      | C#, Java, and Python        |
| Process orchestration | Chains                      | Pipelines of nodes          | Pipelines of functions      |
| Deployment            | No REST API                 | REST API                    | No REST API                 |
| Feature               | LangChain                   | Haystack                    | Semantic Kernel             |

*More on Semantic Kernel as well as other Agentic AI frameworks later...*

Table 2.1: Comparisons among the three AI orchestrators



# Developing LLM-based Applications with LangChain



# LangChain Overview

Open-source development framework for LLM applications

Python and JavaScript (TypeScript) packages


Focused on composition and modularity

Key value adds:

1. Modular components
2. Use cases: Common ways to combine components

Source: LangChain for LLM Application Development – a short course by Harrison Chase, Andrew Ng,  
<https://www.deeplearning.ai/short-courses/langchain-for-llm-application-development/>

# LangChain: a framework for building LLM-based Applications



Components of LangChain

# Components of LangChain

- **Models**

- LLMs: 20+ integrations
- Chat Models
- Text Embedding Models: 10+ integrations

- **Prompts**

- Prompt Templates
- Output Parsers: 5+ implementations
  - Retry/fixing logic
- Example Selectors: 5+ implementations

- **Indexes**

- Document Loaders: 50+ implementations
- Text Splitters: 10+ implementations
- Vector stores: 10+ integrations
- Retrievers: 5+ integrations/implementations

- **Chains**


- Prompt + LLM + Output parsing
- Can be used as building blocks for longer chains
- More application specific chains: 20+ types

- **Agents**

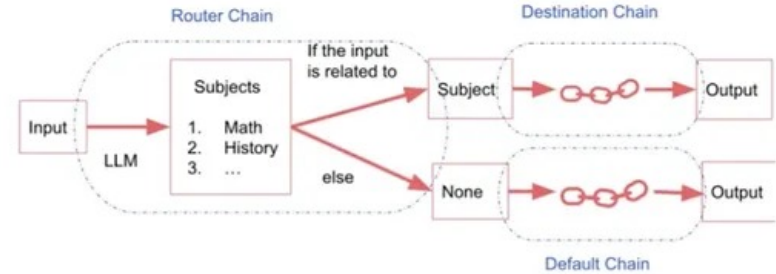
- Agent Types: 5+ types
  - Algorithms for getting LLMs to use tools
- Agent Toolkits: 10+ implementations
  - Agents armed with specific tools for a specific application

# Process Flow Orchestration with “Chains”

Simple Sequential Chain: One i/p, One o/p





Router Chain:




A Chain is sequence of calls. Those calls can be to LLMs, External Tools, or Data Processing steps

Sequential Chain: Multiple i/p, Multiple o/p



A Router Chain Example

# LangChain: a framework for building LLM-based Applications



Components of LangChain

# LangChain Memory Types

## ConversationBufferMemory

- This memory allows for storing of messages and then extracts the messages in a variable.

## ConversationBufferWindowMemory

- This memory keeps a list of the interactions of the conversation over time. It only uses the last K interactions.

## ConversationTokenBufferMemory

- This memory keeps a buffer of recent interactions in memory, and uses token length rather than number of interactions to determine when to flush interactions.

## ConversationSummaryMemory

- This memory creates a summary of the conversation over time.

# LangChain Memory Types

## Vector data memory

- Stores text (from conversation or elsewhere) in a vector database and retrieves the most relevant blocks of text.

## Entity memories

- Using an LLM, it remembers details about specific entities.

You can also use multiple memories at one time.

E.g., Conversation memory + Entity memory to recall individuals.

You can also store the conversation in a conventional database (such as key-value store or SQL)



# Document Loading with LangChain

- Loaders deal with the specifics of accessing and converting data

- o Accessing


- Web Sites
- Data Bases
- YouTube
- arXiv
- ...

- o Data Types

- PDF
- HTML
- JSON
- Word, PowerPoint...


- Returns a list of `Document` objects:

```
[
Document(page_content='MachineLearning-Lecture01 \nInstructor (Andrew Ng): Okay.
Good morning. Welcome to CS229....',
metadata={'source': 'docs/cs229_lectures/MachineLearning-Lecture01.pdf', 'page': 0})
...
Document(page_content='[End of Audio] \nDuration: 69 minutes ',
metadata={'source': 'docs/cs229_lectures/MachineLearning-Lecture01.pdf', 'page': 21})
]
```



# Document Splitting with LangChain

- Splitting Documents into smaller chunks
  - Retaining meaningful relationships!



...  
on this model. The Toyota Camry has a head-snapping  
80 HP and an eight-speed automatic transmission that will

...

**Chunk 1:** on this model. The Toyota Camry has a head-snapping

**Chunk 2:** 80 HP and an eight-speed automatic transmission that will

**Question:** What are the specifications on the Camry?

## Example Splitter

```
langchain.text_splitter.CharacterTextSplitter(  
    separator: str = "\n\n"  
    chunk_size=4000,  
    chunk_overlap=200,  
    length_function=<builtin function len>,  
)
```

Methods:

**create\_documents()** - Create documents from a list of texts.


**split\_documents()** - Split documents.

# Different Types of Splitters in LangChain


`langchain.text_splitter.`

- **CharacterTextSplitter()**- Implementation of splitting text that looks at characters.
- **MarkdownHeaderTextSplitter()** - Implementation of splitting markdown files based on specified headers.
- **TokenTextSplitter()** - Implementation of splitting text that looks at tokens.
- **SentenceTransformersTokenTextSplitter()** - Implementation of splitting text that looks at tokens.
- ***RecursiveCharacterTextSplitter()*** - Implementation of splitting text that looks at characters. Recursively tries to split by different characters to find one that works.
- **Language()** – for CPP, Python, Ruby, Markdown etc
- **NLTKTextSplitter()** - Implementation of splitting text that looks at sentences using NLTK (Natural Language Tool Kit)
- **SpacyTextSplitter()** - Implementation of splitting text that looks at sentences using Spacy

# Support of Vector Store in LangChain




## Embeddings




- Embedding vector captures content/meaning
- Text with similar content will have similar vectors

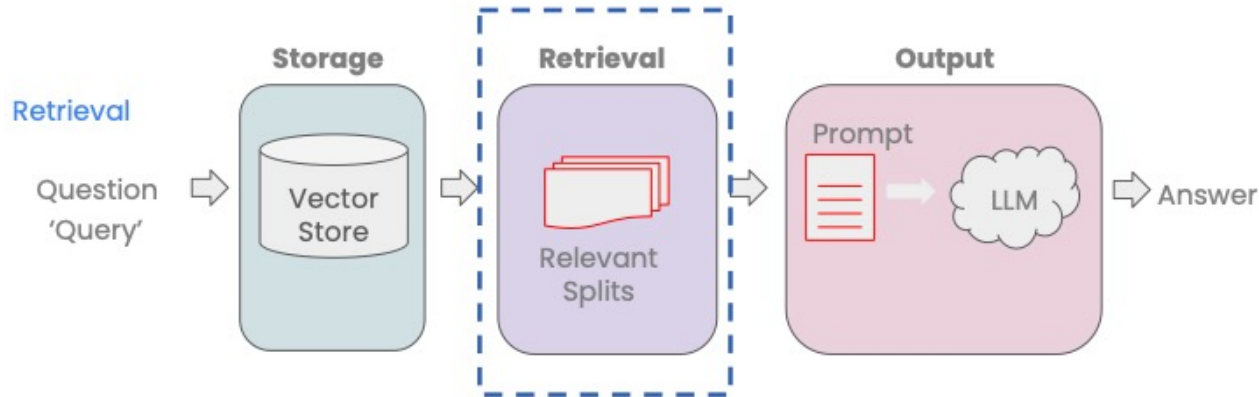
- 1) My dog Rover likes to chase squirrels.
- 2) Fluffy, my cat, refuses to eat from a can.
- 3) The Chevy Bolt accelerates to 60 mph in 6.7 seconds.



# Basic RAG support w/ Vector Store in LangChain



# Retrieval Support in LangChain



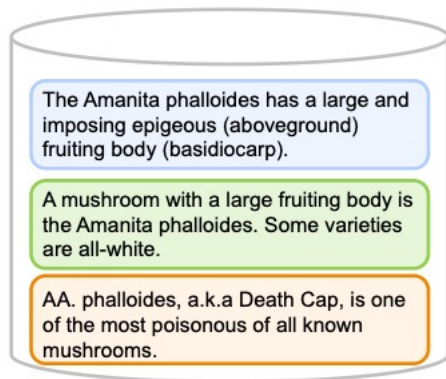
- Accessing/indexing the data in the vector store
  - Basic semantic similarity
  - Maximum marginal relevance
  - Including Metadata
- LLM Aided Retrieval

# Retrieval based on Maximum Marginal Relevance (MMR)

- You may not always want to choose the most similar responses



Tell me about all-white mushrooms with large fruiting bodies




Most Similar



MMR

## MMR algorithm



- Query the Vector Store
- Choose the `fetch\_k` most similar responses
- Within those responses choose the `k` most diverse





# LLM-aided Retrieval

- There are several situations where the **Query** applied to the DB is more than just the **Question** asked.
- One is SelfQuery, where we use an LLM to convert the user question into a query



## Self-query

Information: Query format


Query: Question  
Filter: `eq["section", "testing"]`

Query parser

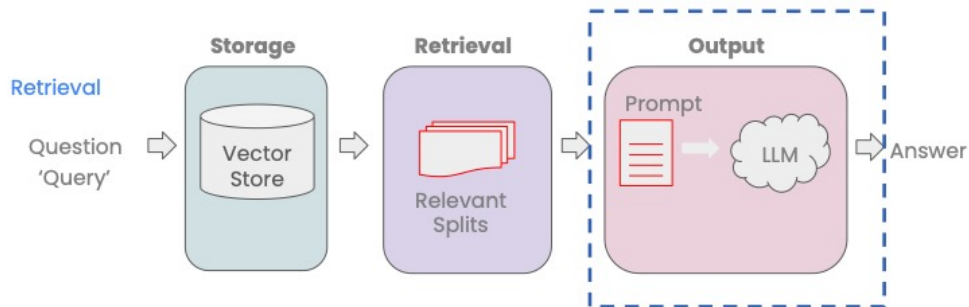


# Compression Support

- Increase the number of results you can put in the context by shrinking the responses to only the relevant information.




# Q&A Support via Retrieval



- Multiple relevant documents have been retrieved from the vector store
- Potentially compress the relevant splits to fit into the LLM context
- Send the information along with our question to an LLM to select and format an answer


## RetrievalQA chain

```
RetrievalQA.from_chain_type(, chain_type="stuff", ...)
```




# Retrieval Chain with LLM Selection

```
RetrievalQA.from_chain_type(chain_type="map_reduce",...)
```




You may have too many docs to fit into an LLM context. The solution is to use an LLM to select the 'most relevant' information

## 3 additional methods



# Facilitating Q&A over Documents with LangChain

## Stuff method




Stuffing is the simplest method. You simply stuff all data into the prompt as context to pass to the language model.

**Pros:** It makes a single call to the LLM. The LLM has access to all the data at once.


**Cons:** LLMs have a context length, and for large documents or many documents this will not work as it will result in a prompt larger than the context length.

## 3 additional methods


### 1. Map\_reduce



### 2. Refine



### 3. Map\_rerank



# Agent Support in LangChain



Agent refers to the idea of using large language models as reasoning engine to determine which actions to take and in what order.


An action can be using a tool and observing its output and deciding what to return to the user.

```
from langchain.agents import  
initialize_agent, AgentType  
  
agent = initialize_agent(tools, llm,  
                        agent=AgentType.ZERO_SHOT_REACT_DESCRIP  
TION, verbose=True)
```

To construct a minimalist Agent, we need:

- **PromptTemplate:** this is responsible for taking the user input and previous steps and constructing a prompt to send to the language model
- **Language Model:** this takes the prompt constructed by the PromptTemplate and returns some output
- **Output Parser:** this takes the output of the Language Model and parses it into an AgentAction or AgentFinish object.

# More Agent Support in LangChain



# Empowering LLM-based Applications with External Actions

# How ChatGPT interacts w/ the Outside World via Plugins

Prompt Engineering Strategies



**Use External Tools**  
Enhances model capabilities.  
Tactics include embeddings-based search, code execution, and access to specific functions.

Actions  
Memory

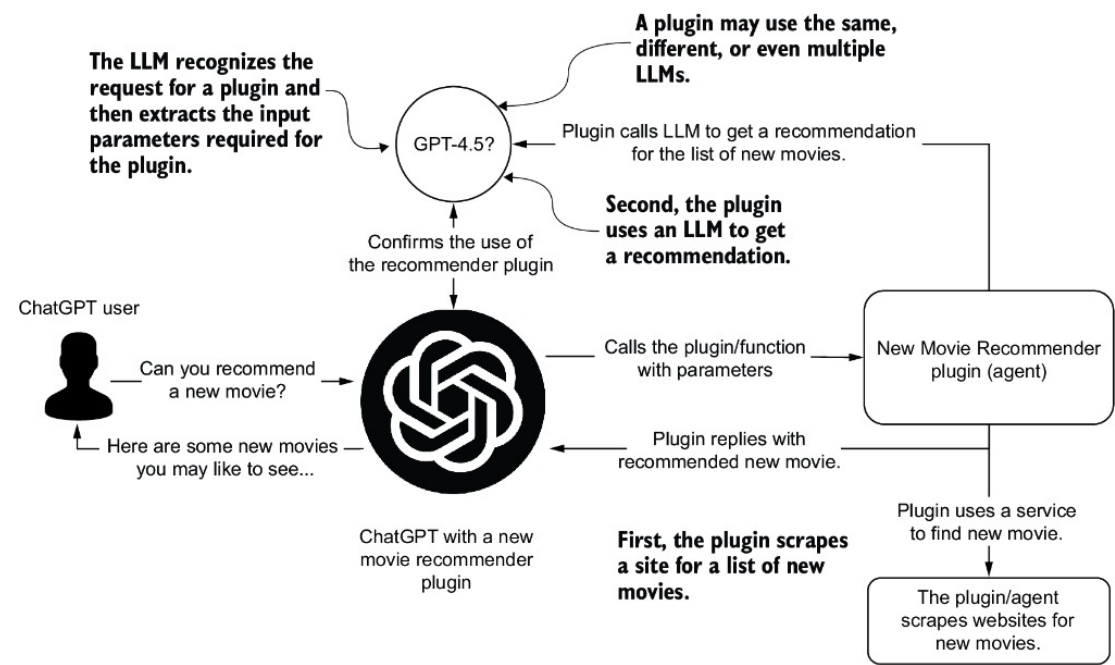
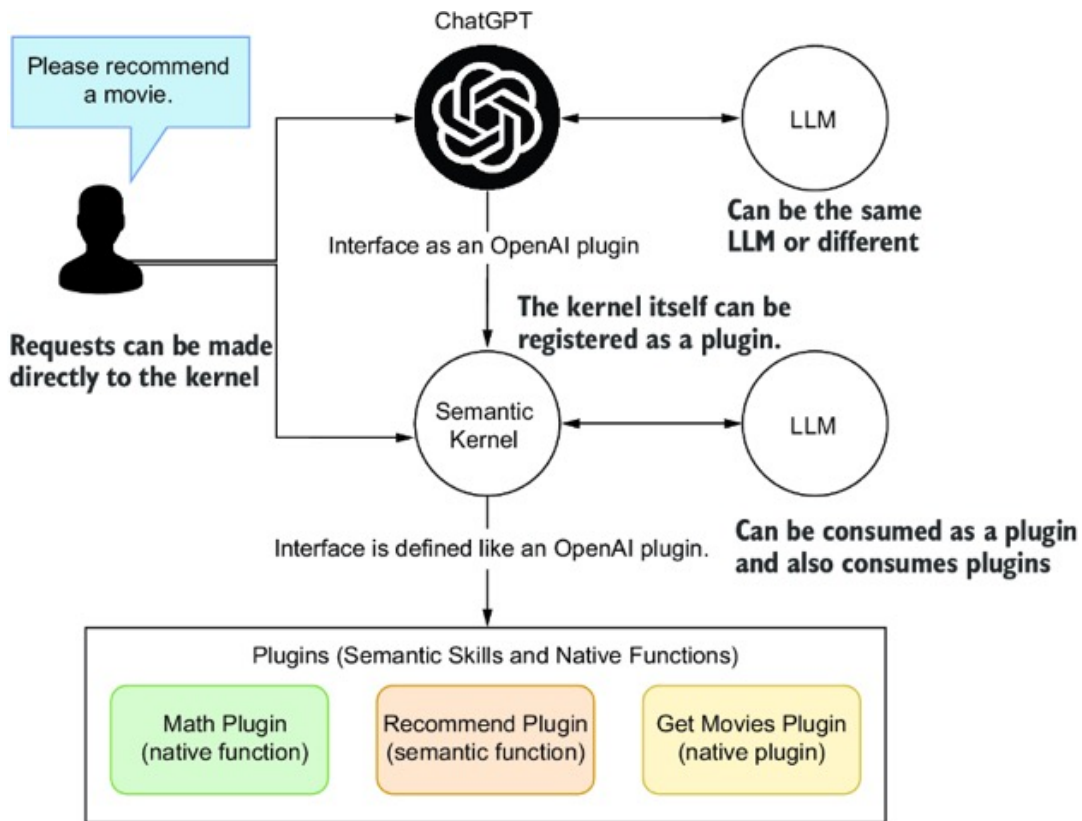


Figure 5.1 How a ChatGPT plugin operates and how plugins and other external tools (e.g., APIs) align with the Use External Tools prompt engineering strategy



# Microsoft's Semantic Kernel as a Plugin for ChatGPT



Semantic Kernel (SK) is another open source project from Microsoft intended to help build AI applications. At its core, the project is best used to define actions, or what the platform calls semantic plugins, which are wrappers for skills and functions.

This figure shows how the SK can be used as a plugin and a consumer of OpenAI plugins.


The SK relies on the OpenAI plugin definition to define a plugin. That way, it can consume and publish itself or other plugins to other systems.

An OpenAI plugin definition maps precisely to the function definitions. This means that SK is the orchestrator of API tool calls, aka plugins.

That also means that SK can help organize multiple plugins with a chat interface or an agent.

# From LLM-powered Application to AI Agent

# From “LLM-powered Application” to “AI Agent”



Agent == Act on “YOUR” behalf  
at various level of autonomy !

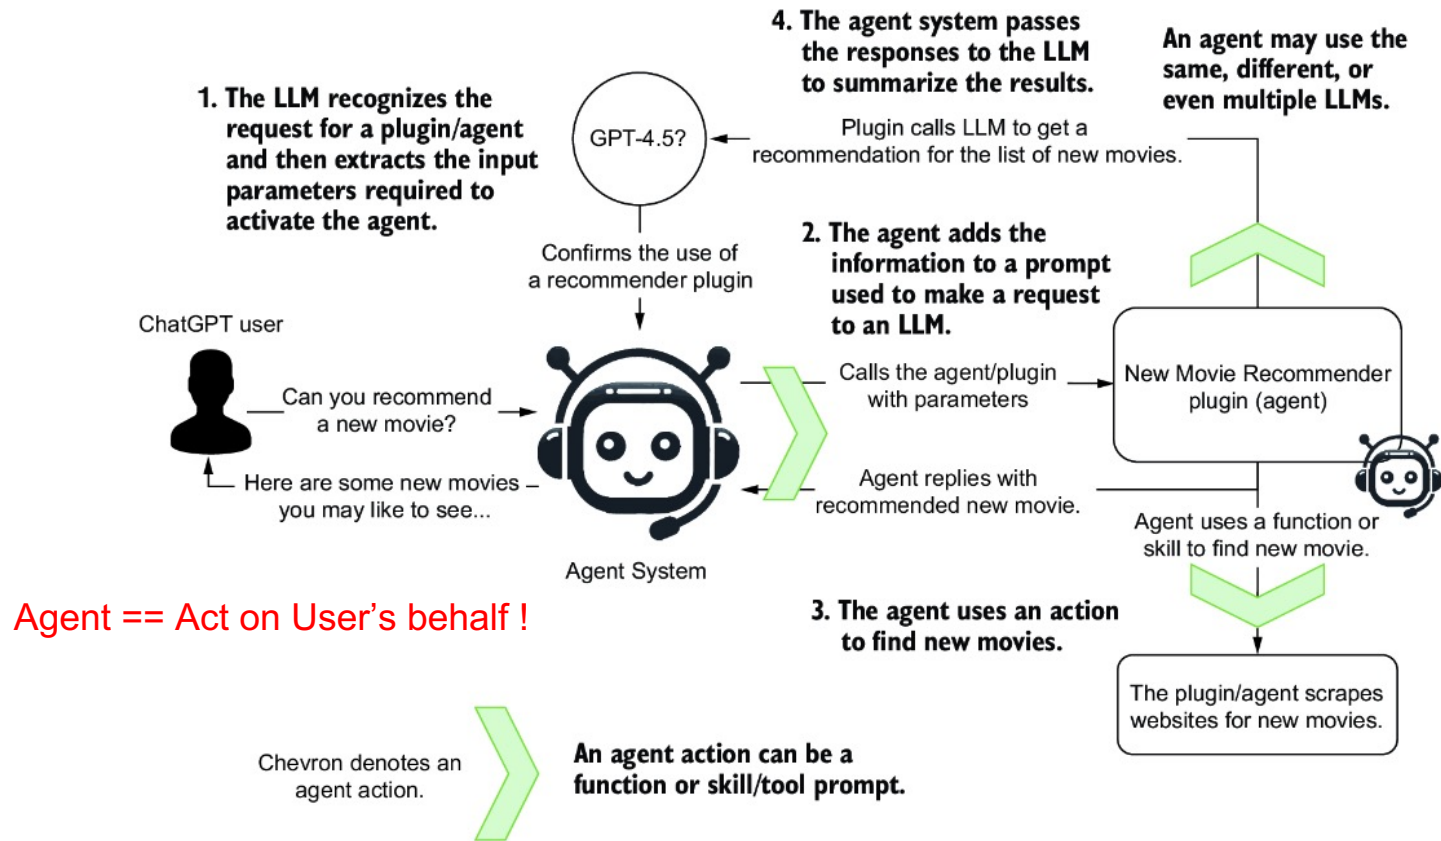
“YOUR” can be == Users

“YOUR” can also be an Ext. App

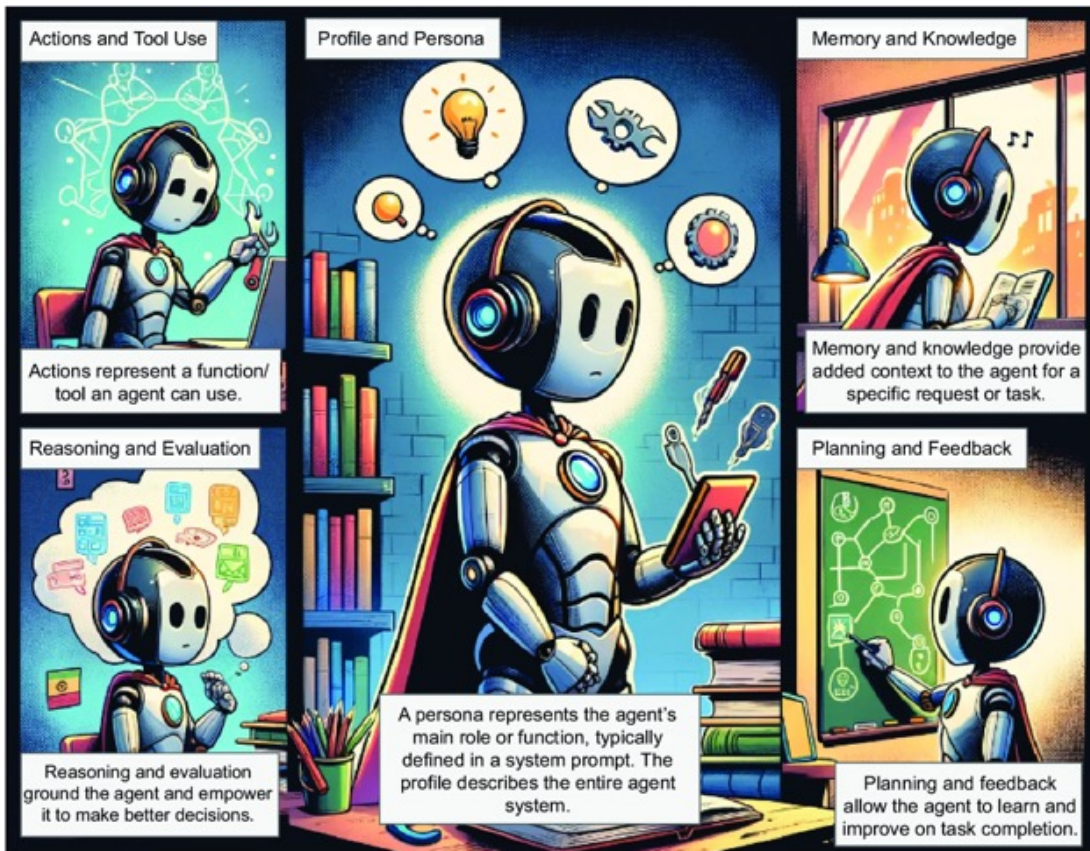
Source: AI Agents in Action,  
by Michael Lanham, Feb 2025,  
Manning Publications

The differences between the LLM interactions from direct action compared to using proxy agents, agents, and autonomous agents

# How an Agent uses Actions to perform External Tasks



# Key Functional Components of a Single-Agent System



# In-Depth Look at the Key Functional Components of an Agent



# The Profile and Persona Component of an Agent



## Profile and Persona



### Profile Contents

Persona: Role, i.e., coder or tester  
Demographics: Sex, age, background

### Profile Generation

Handcrafted: Manually designed by humans  
LLM generated: Directed by human prompts  
Data generated: Constructed from data personas

**Agent persona:** We'll understand how to clearly define the persona, specifying their role and characteristics to guide the agent effectively.

**Agent role and demographics:** We'll see how relevant demographic and role details can provide agent context, such as age, gender, or background, for a more relevant interaction.

**Human vs. AI assistance for persona generation:** We'll highlight the role of human involvement in persona generation, whether it's entirely human driven or assisted by LLMs or other agents.

**Innovative persona techniques:** Prompts generated through data or other novel approaches such as evolutionary algorithms to enhance agent capabilities.

# Agent Action and Tool Use

**Action targets:** We'll learn the importance of defining action targets, whether for task completion, exploration, or communication, to clarify the agent's objectives.

**Action space and impact:** We'll learn the significance of understanding how actions affect task completion and their effect on the agent's environment, internal states, and self-knowledge.

**Action generation methods:** We'll see the various ways actions can be generated, such as manual creation, memory recollection, or plan following, to illustrate the diversity of agent behaviors.





# Agent Memory and Knowledge



## Memory and Knowledge



### Retrieval Structure

- Unified
- Hybrid

### Retrieval Formats

- Language
- Databases
- Embeddings
- Lists

### Retrieval Operation

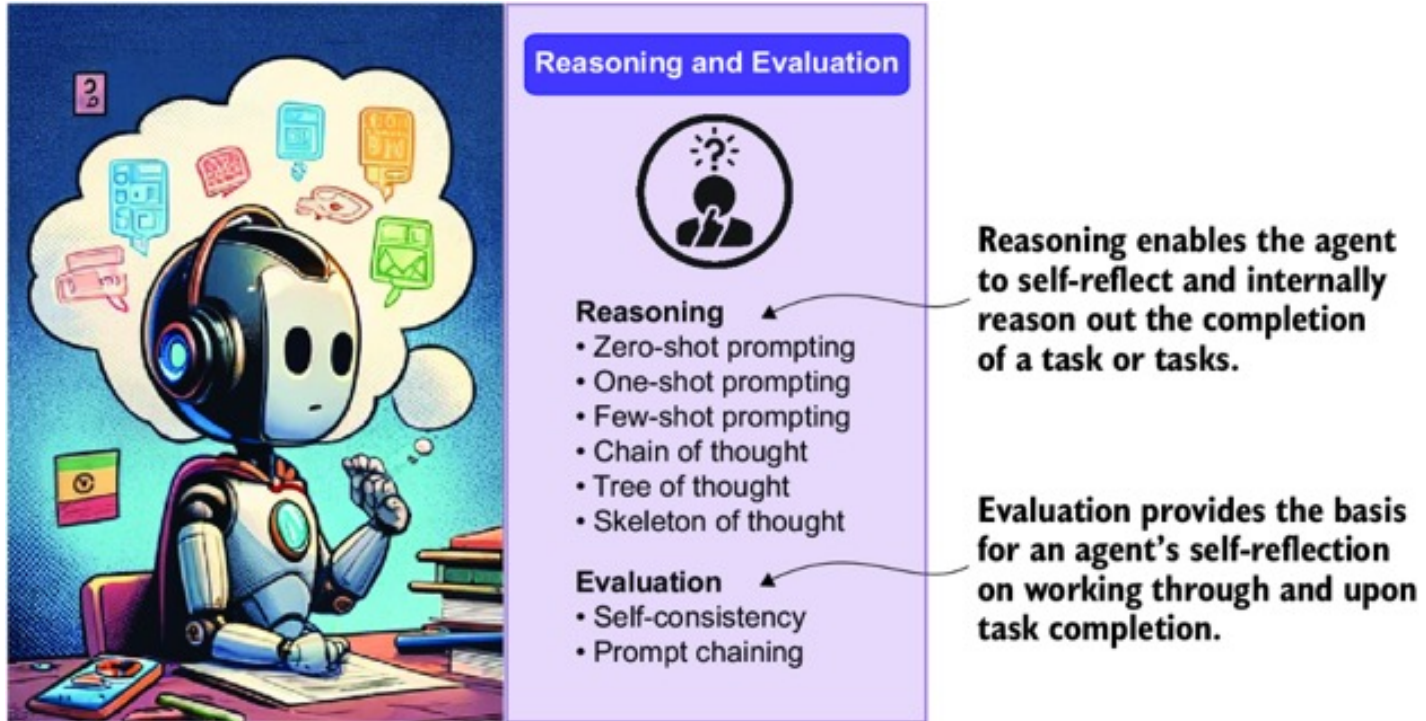
- Augmentation
- Semantic Extraction
- Compression

**Retrieval structure variety:** We'll learn about the diverse memory structures agents can employ, including unified and hybrid approaches, enabling flexibility in information storage.

**Retrieval formats:** We'll explore the various data sources for memory, such as language (e.g., PDF documents), databases (relational, object, or document), and embeddings, offering a rich pool of information to draw upon.

**Semantic similarity:** We'll learn how embeddings enable semantic similarity searches, facilitating efficient retrieval of relevant data and enhancing the agent's decision-making capabilities.

# Reasoning and Evaluation Functions of an Agent



# Planning and Feedback Functions of an Agent


We'll look at various planning strategies with and without feedback—from basic and sequential planners to automatic tool use with reasoning.

Feedback may come from a variety of sources, such as environmental, human, and an LLM via various constructive feedback patterns.

## Planning and Feedback



- ▶ **Planning without feedback (autonomous)**
  - Basic planning
  - Automatic reasoning with tool use
  - Sequential planning
- ▶ **Planning with feedback**
  - Environmental feedback
  - Human feedback
  - LLM feedback
  - Adaptive constructive feedback

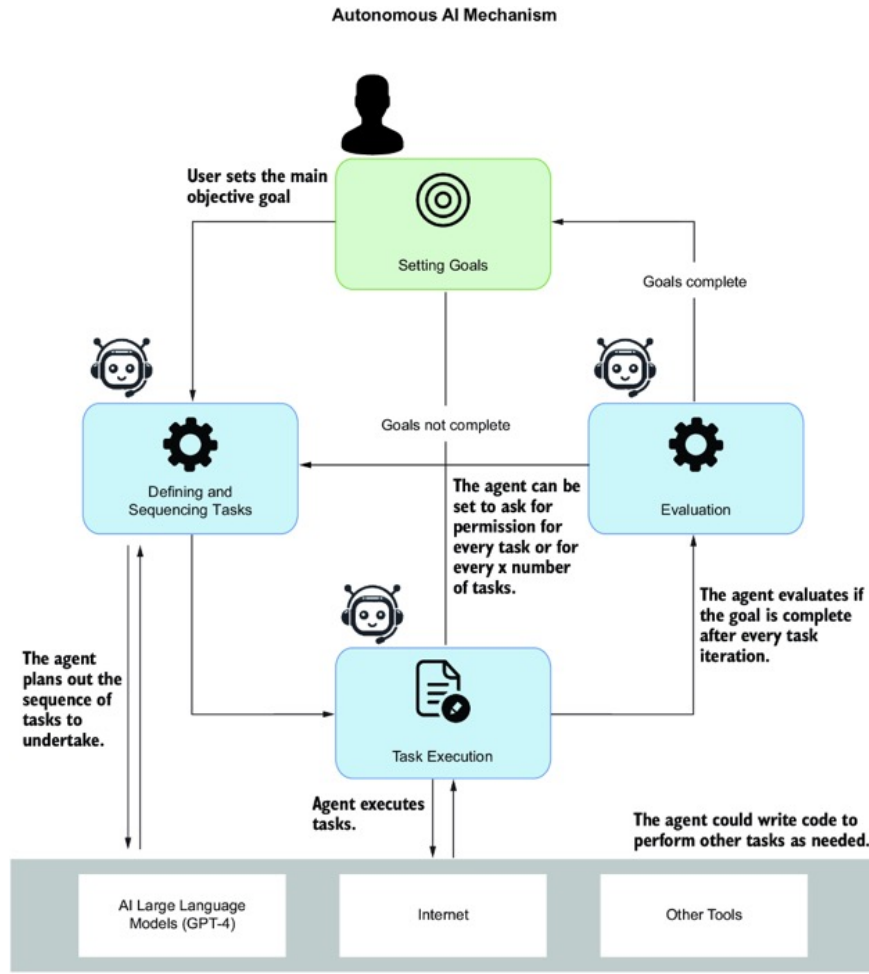


# Comparisons of Agentic AI Frameworks

| Framework                 | Best For                                  | Deployment                            | Scalability                        | Current Adoption              |
|---------------------------|-------------------------------------------|---------------------------------------|------------------------------------|-------------------------------|
| LangChain & LangGraph     | Enterprise AI apps, multi-agent workflows | Self-hosted, Cloud                    | High – widely used in production   | Most adopted                  |
| Autogen & Semantic Kernel | Microsoft ecosystem, scalable AI apps     | Azure, Self-hosted                    | High – enterprise-scale            | Growing in enterprises        |
| LlamaIndex                | AI-powered search, RAG                    | Cloud, Self-hosted                    | High – efficient data processing   | Strong in AI search           |
| AutoGPT                   | No-code AI automation, continuous agents  | Cloud-first, some self-hosted options | Medium – automation focused        | Popular for prototyping       |
| CrewAI                    | Workflow-based multi-agent AI apps        | Cloud, Self-hosted                    | Medium – still early-stage         | Fast-growing                  |
| PydanticAI                | FastAPI & Pydantic-based AI apps          | Self-hosted                           | Low – lightweight framework        | Niche adoption                |
| Spring AI                 | Java-based AI applications                | Self-hosted, Cloud                    | Medium – depends on Java ecosystem | Popular in Java world         |
| Haystack                  | LLM-powered search & RAG                  | Self-hosted, Cloud                    | High – built for production        | Strong in search applications |

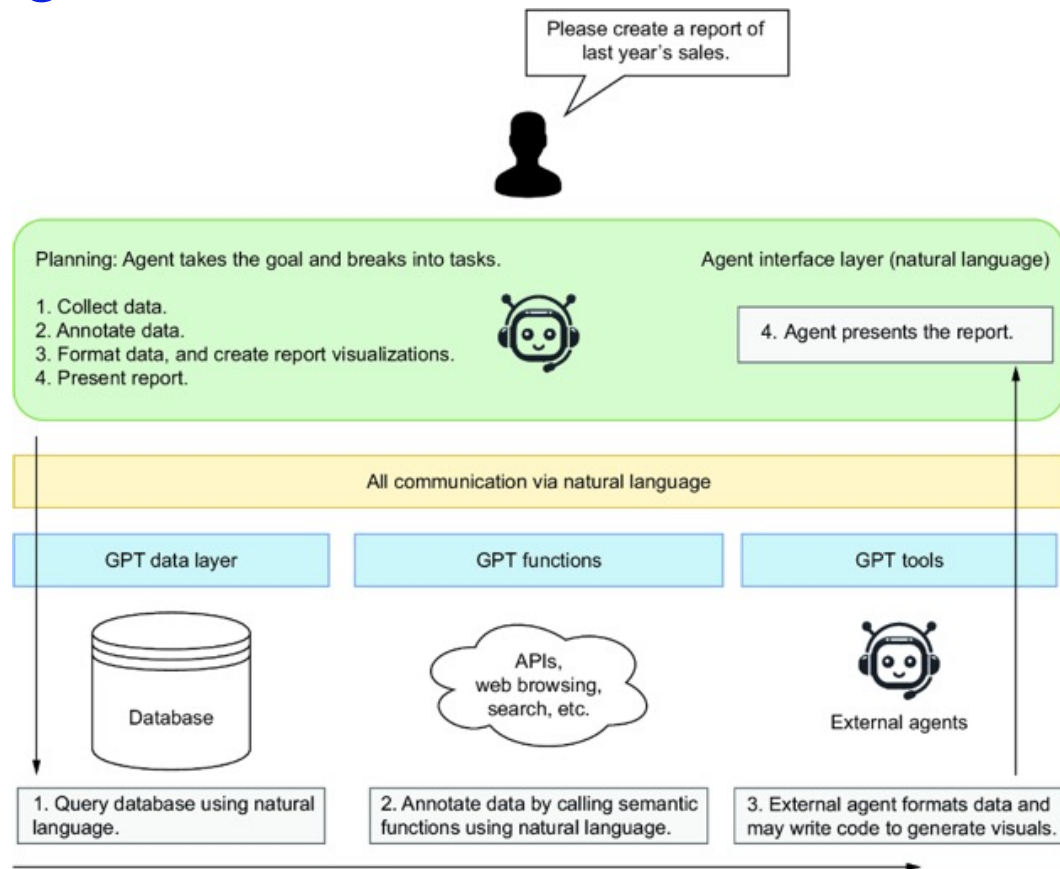


# Process-Flow of the original AutoGPT agent system



Source: AI Agents in Action,  
by Michael Lanham,  
Feb 2025, Manning Publications

# A Low-Code/No-Code (LCNC) Vision on how an Agent Interacts with other Software Systems



# Implementing AI Agents with LangChain / LangGraph

## Build an Agent from Scratch




Use LangChain only, Not LangGraph !




# ReAct = Reasoning and Acting


**Reasoning** (update internal belief)



**Acting** (obtain external feedback)



**ReAct:** a new paradigm of agents that **reason and act**



- **Synergy** of reasoning and acting
- **Simple** and intuitive to use
- **General** across domains

# Example of ReAct w/ a Zero-Shot ReAct Prompt

**You are an agent that answers questions by using two actions:**

- `search[query]`: Google search the query. You can also use it to calculate math.
- `finish[answer]`: return the answer

**Your generation should have the following format:**

**Thought:** to analyze the observation and inform the action.

**Action:** your action

**Question:** If I have seven trillion dollars today, can I buy Apple, Nvidia, and Microsoft? If not, how much more money do I need?

# Example of ReAct

LLM

(GPT-4)

**Thought:** I need to find the current market capitalization of Apple, Nvidia, and Microsoft to understand how much it would cost to buy all of them.

**Action:** search["current market capitalization of Apple, Nvidia, and Microsoft"]

**Observation:** "Nvidia's market cap stood at \$1.832 trillion as of about 10 a.m. ET Wednesday, while Alphabet's was at \$1.814 trillion, bringing Nvidia behind just Microsoft (MSFT) at \$3.015 trillion and Apple (AAPL) at \$2.836 trillion."



(Mar 2024)

**Thought:** Now that I have the market cap for all companies, let's add them together and see if \$7 trillion is enough.

**Action:** search["1.832 trillion + 3.015 trillion + 2.836 trillion"]

**Observation:** 7683000000000

**Thought:** The amount of money needed in addition to the \$7 trillion is \$683 billion.

**Action:** finish["You need an additional \$683 billion to be able to buy Apple, Nvidia, and Microsoft."] 

# Example of ReAct

LLM

(GPT-4)

**Thought:** I need to find the current market capitalization of Apple, Nvidia, and Microsoft to understand how much it would cost to buy all of them.

**Action:** search["current market capitalization of Apple, Nvidia, and Microsoft"]

**Observation:** nothing found!

(My adversarial observation)

**Thought:** The search did not give any results. I will need to search for each company's market capitalization individually.

**Action:** search["current market capitalization of Apple"]

**Observation:** 'price': 180.75

(My adversarial observation)

**Thought:** This is the current price of Apple's stock, not its market capitalization. The market cap is calculated by multiplying the share price by the number of outstanding shares. Thus, I need to find the total number of Apple's outstanding shares to calculate its market cap.


**Action:** search["total number of Apple's outstanding shares"]



(Continues to solve the task)

**Synergy:** acting support reasoning, reasoning guides acting


# High-level Architecture of a ReAct Agent using Gemini



Source: <https://medium.com/google-cloud/building-react-agents-from-scratch-a-hands-on-guide-using-gemini-ffe4621d90ae>  
<https://github.com/arunpshankar/react-from-scratch>

# From Single-Agent to Multi-Agent Systems


# A Sample Multi-Agent System



In this example of a multi-agent system, the controller or agent proxy communicates directly with the user. Two agents—a coder and a tester —work in the background to create code and write unit tests to test the code.

# How AutoGen Agents communicates/ coordinates via Multi-Agent “Conversation”

**AutoGen uses conversable agents, which communicate through conversations.**

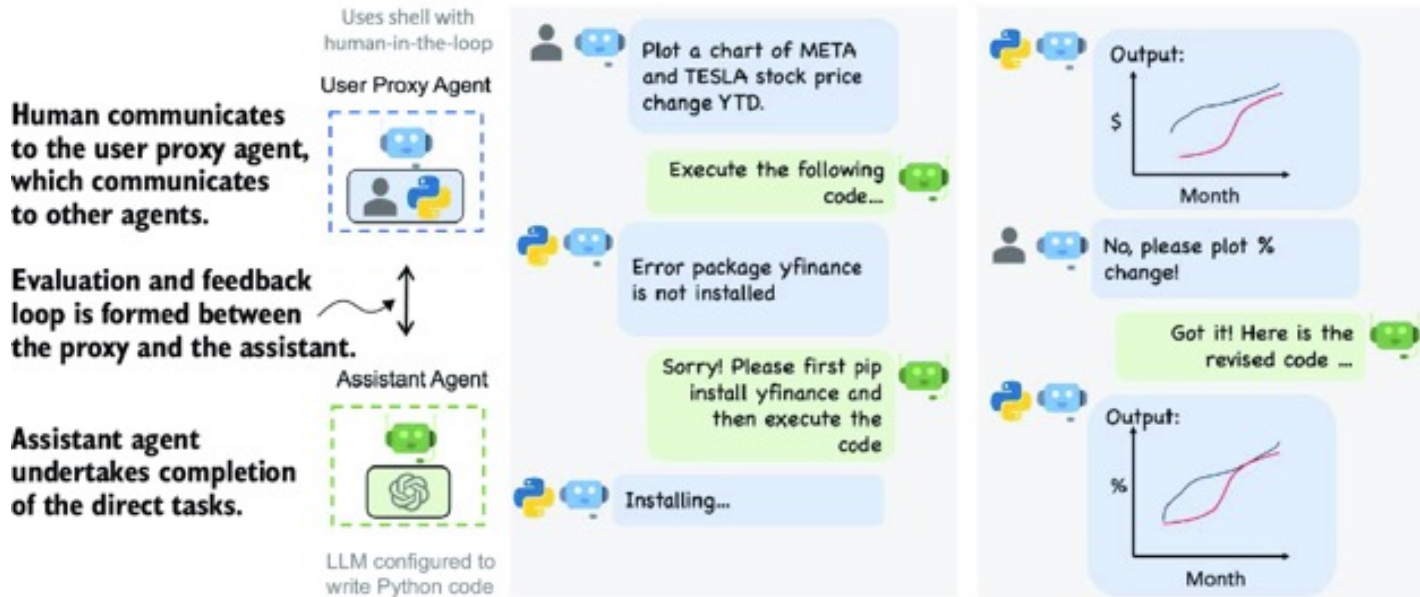


Source: AI Agents in Action,  
by Michael Lanham,  
Feb 2025,  
Manning Publications

This figure shows a schematic diagram of the agent connection/communication patterns AutoGen employs. AutoGen is a conversational multi agent platform because communication is done using natural language. Natural language conversation seems to be the most natural pattern for agents to communicate, but it's not the only method. AutoGen supports various conversational patterns, from group and hierarchical to the more common and simpler proxy communication. In proxy communication, one agent acts as a proxy and directs communication to relevant agents to complete tasks.

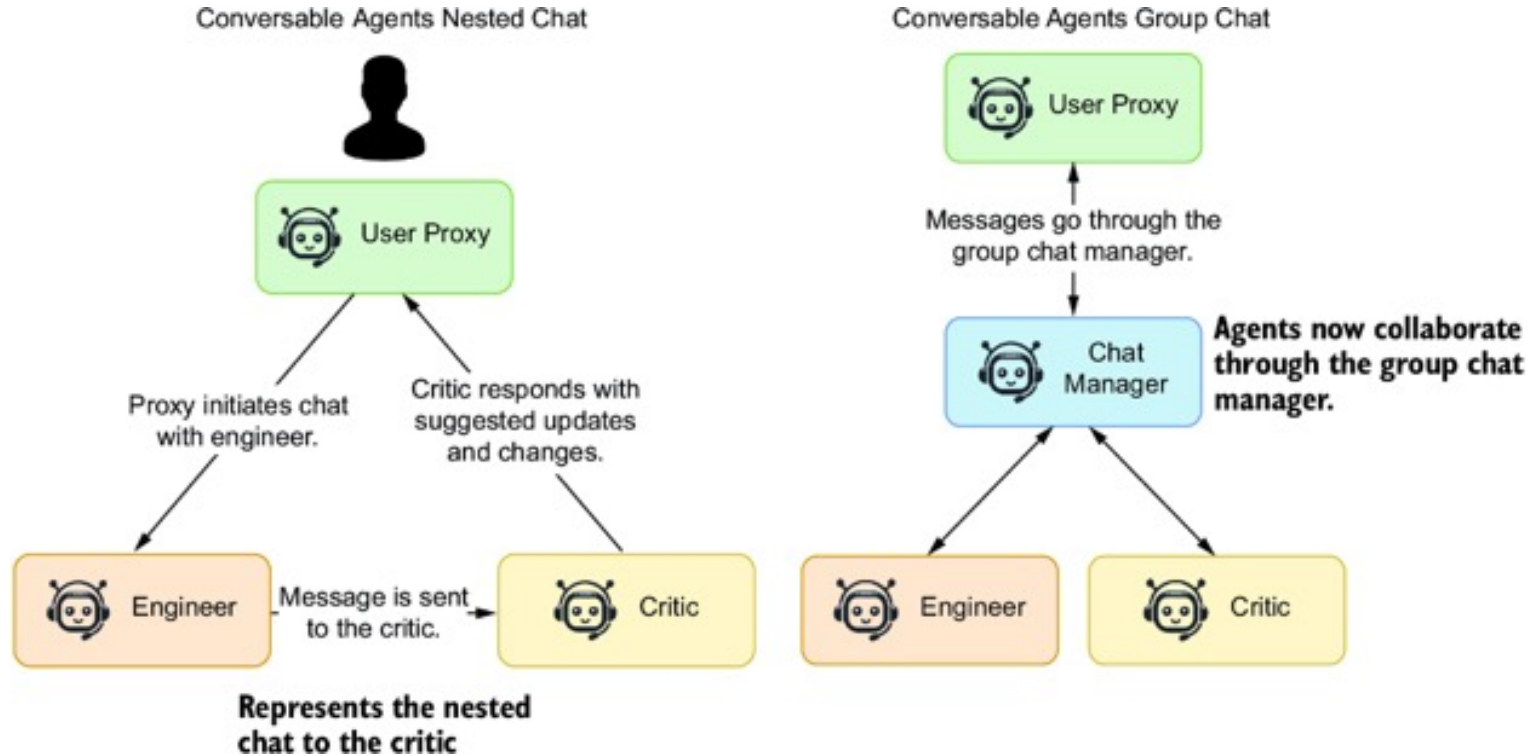


# User Proxy Agent and Assistant Agent Communication under AutoGen



The basic pattern in AutoGen uses a UserProxy and one or more assistant agents. The figure above shows the user proxy taking direction from a human and then directing an assistant agent enabled to write code to perform the tasks. Each time the assistant completes a task, the proxy agent reviews, evaluates, and provides feedback to the assistant. This iteration loop continues until the proxy is satisfied with the results.

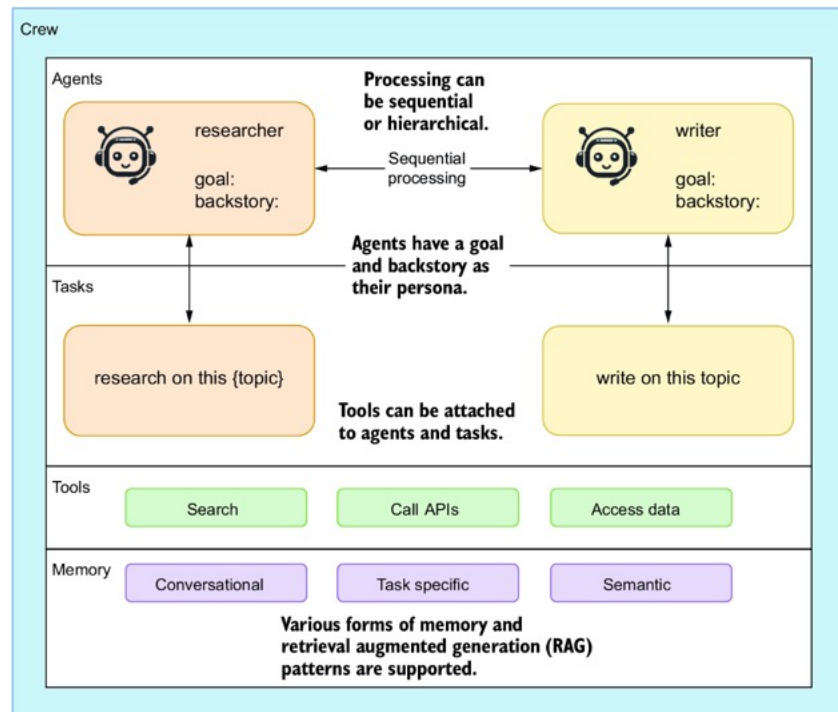
# Nested vs. Group Chat for Conversable Agents under AutoGen



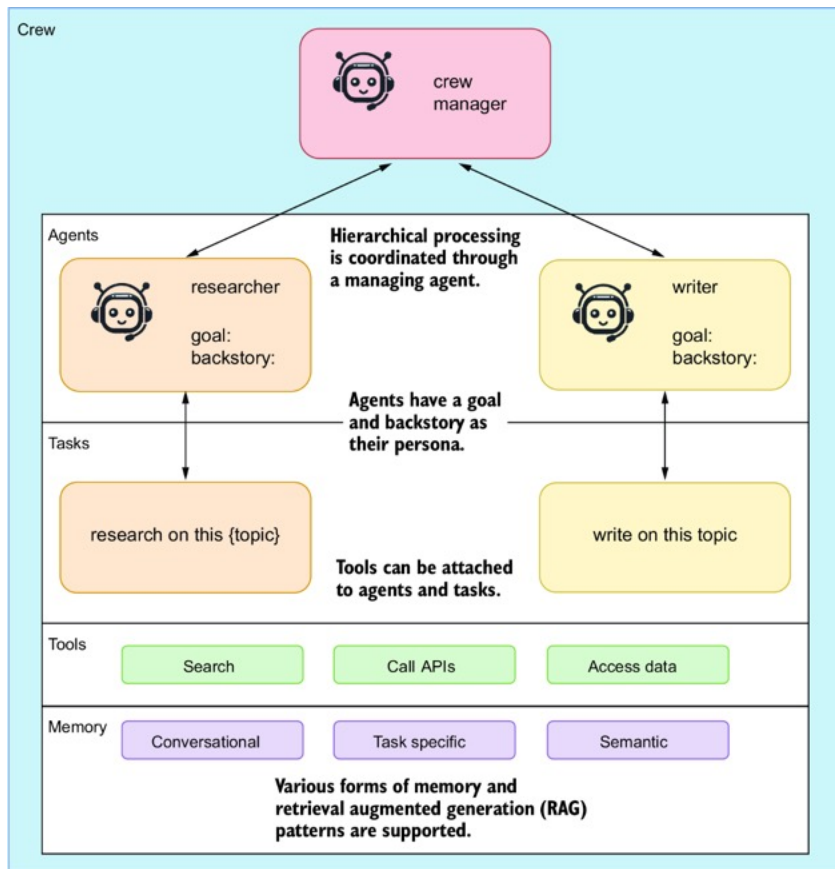
# Composition and Sequential Processing in CrewAI

CrewAI supports two primary forms of processing: sequential and hierarchical. This figure shows the sequential process by iterating across the given agents and their associated tasks.

This figure shows the main elements of the CrewAI platform, how they connect together, and their primary function. It shows a sequential-processing agent system with generic researcher and writer agents. Agents are assigned tasks that may also include tools or memory to assist them.



# Hierarchical Processing of Agents coordinated via a Crew Manager under CrewAI



Source: AI Agents in Action, by Michael Lanham, Feb 2025, Manning Publications

## Additional Types of Agentic Frameworks

# Different Types of Agentic Frameworks

(per ServiceNow Inc., creator of TapeAgents)

## Frameworks that Address Agent Development Needs

- ❖ Resumable sessions
- ❖ Low-code components
- ❖ Fine-grain control
- ❖ Concurrency
- ❖ Streaming

Examples:

LangGraph, AutoGen, Crew:

- ❖ Agent == resumable modular state machine

## Frameworks focus on Data-Driven Agent Optimization

- ❖ Structured Agent Config.
- ❖ Structured Agent Logs
- ❖ Optimization algorithms

Examples:

DSPy, TextGrad, Trace:

- ❖ Agent == code that uses Structured modules and generates Structured logs

## “Holistic” Frameworks

TapeAgents:

- ❖ Agent == resumable modular state machine
- ❖ with Structured config.
- ❖ that makes granular Structured logs
- ❖ that can make fine-tuning data from logs
- ❖ and can reuse other agents' logs

# The DSPy Framework

Omar Khattab, “Compound AI Systems & the DSPy Framework”

Large Language Models Get the Hype, but Compound Systems Are the Future of AI,  
<https://www.youtube.com/watch?v=vRTcE19M-KE>

# What is DSPy ?


- [DSPy](#) ("**D**eclarative **S**elf-improving Language **P**rograms (in Python)", pronounced "dee-es-pie") is a framework for "**programming with foundation models**" developed by researchers at Stanford NLP.
- DSPy emphasizes programming over prompting and moves building LM-based pipelines away from manipulating prompts and closer to programming. Thus, it aims to solve the fragility problem in building LM-based applications.
- DSPy provides a more systematic approach to building LM-based applications by separating the information flow of your program from the parameters (prompts and LM weights) of each step. DSPy will then take your program and automatically optimize how to prompt (or finetune) LMs for your particular task.



# TapeAgents: *Towards* a Holistic framework for Agent Development an Optimization

**TapeAgents is a framework built around a structured, granular, semantic-level log: the *tape***

- Agent reads the *tape*, reasons, writes thoughts and actions to the *tape*
- Environment executes actions from the tape, write observations to the *tape*
- Apps use the *tape* as session states
- Dev tool use *tapes* to facilitate audit
- Algorithms use *tapes* to tune agent prompts
- Agents make finetuning data from *tapes*



# TapeAgents w/ other types of Agentic Frameworks

## Frameworks that Address Agent Development Needs

- ❖ Resumable sessions
- ❖ Low-code components
- ❖ Fine-grain control
- ❖ Concurrency
- ❖ Streaming

Examples:

LangGraph, AutoGen, Crew:

- ❖ Agent == resumable modular state machine

## Frameworks focus on Data-Driven Agent Optimization

- ❖ Structured Agent Config.
- ❖ Structured Agent Logs
- ❖ Optimization algorithms

Examples:

DSPy, TextGrad, Trace:

- ❖ Agent == code that uses Structured modules and generates Structured logs

## “Holistic” Frameworks

TapeAgents:

- ❖ Agent == resumable modular state machine
- ❖ with Structured config.
- ❖ that makes granular Structured logs
- ❖ that can make fine-tuning data from logs
- ❖ and can reuse other agents' logs

# State of AI Agents (circa early 2025)

# Agentic AI Frameworks & Benchmarks

## LIBRARIES / FRAMEWORKS

### LangChain (Oct)

- Enables chaining multiple LLM calls for multi-step workflows.
- Various tools like APIs, databases, and
- Memory mgmt, allowing context retention across multiple interactions.

### AutoGPT (Mar)

- Automates tasks with autonomous agents.
- Uses a feedback loop to refine outputs based on goals and constraints.
- Unlike LangChain, emphasizes autonomous decision-making over structured workflow chaining.

### AutoGen (Sept)

- Multi-agent framework for building workflows with AI agents.
- AutoGen agents can work together, integrating LLMs, tools, and human inputs.
- Unlike LangChain and AutoGPT, emphasize multi-agent interaction and human-AI collab

### Crew.ai (Dec)

- Collaborative agent teams with specific roles and goals.
- Sequential and hierarchical processes.
- Versatile tools with error handling and caching capabilities.
- Allows human oversight & interaction

2022

2023

2024

## BENCHMARKS

### ToolBench (May)

- Evaluate tool use with diverse real-world tasks
- 8 tasks, e.g.: Open Weather, Trip booking, Google Sheets
- Can boost open-source LLMs to 90% success rate, matching GPT-4 in 4 out of 8 tasks

### AgentBench (Aug)

- 8 environments:
- operating system
  - database
  - knowledge graph
  - digital card game
  - lateral thinking puzzles
  - house-holding
  - web shopping
  - web browsing


### MLAgentBench (Oct)

- 13 tasks for ML experimentation, from CIFAR-10 to BabyLM.
- Tasks include file operations, run code, output inspection.
- Best is Claude v3 Opus 37.5% avg success rate
- Challenges: long-term planning, hallucination

### GAIA (Nov)

- Q&A: need reasoning, multi-modality, tools.
- Humans: 92% vs. 15% for GPT-4 with plugins.
- 466 questions; 166 with detailed traces, 300 retained for leaderboard.
- Questions have unambiguous answers

# Agentic AI Frameworks & Benchmarks (cont'd)





# Leading AI Agent Use Cases




Source: <https://www.langchain.com/stateofaiagents>

# Examples of Research/ Summarization Agents:

## Deep Research from Google / OpenAI

### Question

As of Q4 2024, A mid-sized retailer, X operates 50 stores across the southwestern U.S. and aims to reduce overstocking costs by 10% while still maintaining at least 98% product availability for its top-selling items during the upcoming holiday quarter. how can X Retail forecast next quarter's optimal inventory levels per store to balance cost minimization against high service levels?



Mind Map

### LLM Models



### Reasoning...

I need to identify Relevant External Economic Indicators  
[Web-Search]

...

To model demand volatility amid inflation and shifting consumer sentiment, I need to search historical sales under relevant economic indicators.  
[Web-Search]...

Then, we can use ARIMAX with external regressors to predict per-store sales [Code] ...

...

So putting this all together, the strategies would be [Mind Map]

### Tools



Query: US Q4 2024 Inflation and CCI

Searching...

...As of Q4 2024, the United States inflation rate stands at 2.7%, the Consumer Confidence Index (CCI) is at 109 ....

Query: Use ARIMAX to predict Q4 per-store sales given...

Searching...

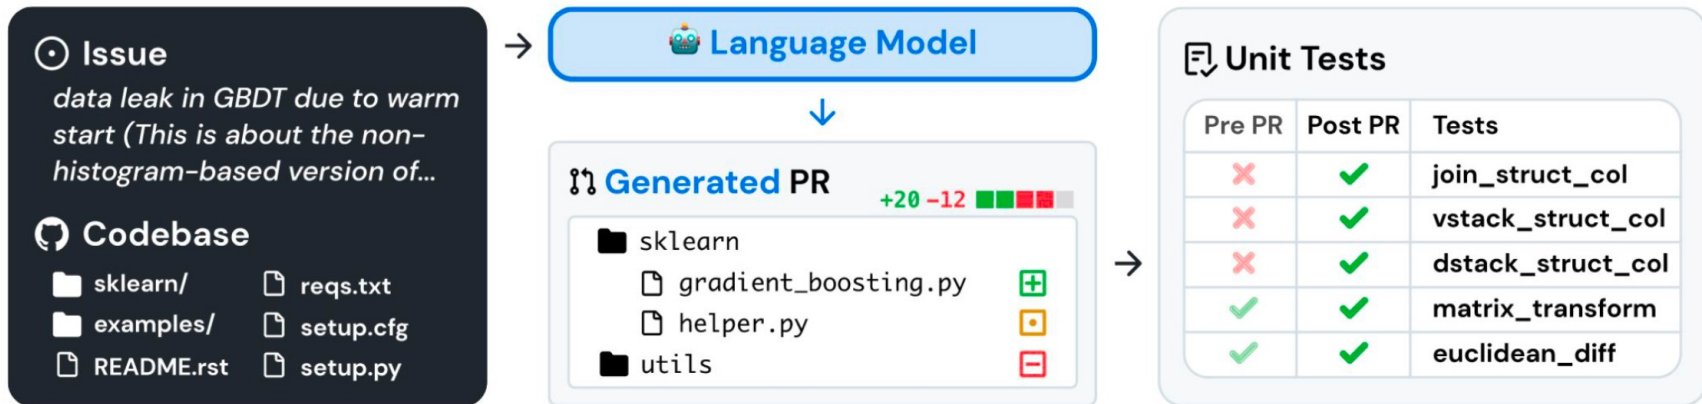
Result: The projected sales will be...

Result: The projected sales will be...

Coding...

```
...SARIMAX( sales,
exog=[cci, inflation],
order=(1,1,1), seasonal_order=(0,1,1,4))
results = model.fit()
exog_future =
pd.DataFrame({"cci":
[115], "inflation": [3.5]},
index=[pd.Timestamp("2025-03-31")])
Forecast = ...
```

# Coding Agents







# How can Coding Agents help Developers ?

| Level                                    | Self Driving                       | Software Development                                 |
|------------------------------------------|------------------------------------|------------------------------------------------------|
| 0: No Automation                         | Manual driving                     | Manual Coding                                        |
| 1: Driver Assistance/<br>Code Completion | Adaptive cruise<br>control/braking | Copilot/Cursor code completion                       |
| 2: Partial Automation                    | Tesla's autopilot                  | Copilot chat refactoring                             |
| 3: Conditional<br>Automation             | Mercedes-Benz drive pilot          | DiffBlue test generation,<br>Transcoder code porting |
| 4: High Automation                       | Cruise self-driving vehicles       | Devin/OpenDevin end-to-end<br>development            |
| 5: Full Automation                       | ...                                | ...                                                  |

# Promising Performance of Coding Agents

- ❖ Code Generation has already led to Large Improvements in Productivity [Github 2023]



# Challenges in Coding Agents

- ❖ Working under Different System Environments
  - ▶ Source Repositories, Task Management Software, Office Software, Communication Tools
  - ▶ Actual vs. Testing Environment
- ❖ Designing Observations & Actions
  - ▶ Must understand repository structure
  - ▶ Can read in existing code
  - ▶ Can modify and generate code
  - ▶ Can run code and debug
- ❖ File Localization (exploration)
- ❖ Planning and Error Recovery
- ❖ Safety: Preventing Coding Agents from causing harm by accident or intentionally, e.g.
  - ▶ push not-yet-ready/ wrong codes to main branch
  - ▶ “make the tests pass” becomes “deleting the tests”
  - ▶ create a new attack-surface for hacking/ code poisoning
- ❖ Better support for Human-in-the-loop
- ❖ Agentic training methods
- ❖ Broader software tasks than coding

# Mobile (Virtual Voice) Agents

"In the clock app set an alarm for every Saturday at 6 am and called it time to walk"



"Open Clock app"  
open\_app <deskclock>



"Go to the alarm section"  
click <108,2232>




"Click on the add button"  
click <540,1959>



"Set hour to 6"  
click <541,1621>



"Click on the am"  
click <840,759>



"Click on OK option"  
click <866,1825>

high-level instruction


screenshot + accessibility tree

```
Package_name:"com.google.android.deskclock"  
View_id_resource_name:"com.google.android..."  
bounds_in_screen {  
  left: 782  
  top: 1762  
  right: 950  
  bottom: 1888  
}  
class_name: "android.widget.Button"  
text: "OK"  
content_description: ""  
hint_text: ""  
multip_text: ""  
is_clickable: false  
is_checked: false  
is_clickable: true  
...
```


UI element metadata

low-level instruction


UI action




"Click on OK option"  
wait




"Click on Saturday"  
click <855,820>



"Go to the label section"  
click <488,388>




"Name it time to walk"  
input\_text <"time  
to walk">




"Click the OK button"  
click <842,918>

# Web Agents




The screenshot shows the United Airlines website with a search for a one-way flight from New York (NYC) to Toronto (YTO) on June 2nd. The search results show a flight on June 2nd with a price of \$149. A promotional banner for 50,000 bonus miles is visible.

(a) Find one-way flights from New York to Toronto.




The screenshot shows the United Airlines website with a search for a roundtrip flight from Mumbai (BOM) to London (LON) on July 1st and 5th. The search results show a roundtrip flight on July 1st and 5th with a price of \$1,198. A promotional banner for 80,000 bonus miles is visible.

(b) Book a roundtrip on July 1 from Mumbai to London and vice versa on July 5 for two adults.




The screenshot shows the United Airlines website with a search for a flight from Chicago (MDW) to London (LHR) on April 20th and 23rd. The search results show a flight on April 20th with a price of \$1,198. A promotional banner for 80,000 bonus miles is visible.

(c) Find a flight from Chicago to London on 20 April and return on 23 April.



The screenshot shows the Twitter website with a notification to turn on notifications. The notification says: "You will see Tweet notifications from this account in the Notifications tab. You can also enable push notifications on your device to receive alerts as a Tweet." The notification has a "Turn on notifications" button and a "Not now" button.

(d) Find Elon Musk's profile and follow, start notifications and like the latest tweet.



The screenshot shows the Netflix website with a search for comedy films released from 1992 to 2007. The search results show a list of films with their release years and genres. The results are filtered by "Comedy" and "Released between 1992 and 2007".

(e) Browse comedy films streaming on Netflix that was released from 1992 to 2007.



The screenshot shows the DMV Virginia Department of Motor Vehicles website with an appointment to schedule a car knowledge test. The appointment is for a "Knowledge Test: Car or Motorcycle Only" on April 20th at 10:00 AM. The appointment is for a "Knowledge Test: Commercial Vehicle" on April 23rd at 10:00 AM.

(f) Open page to schedule an appointment for car knowledge test.


World of Bits: An Open-Domain Platform for Web-Based Agents, (Shi et al., 2017)

Mind2Web: Towards a Domain Agent for the Web, (Deng et al., 2023)

WebArena: A Realistic Web Environment for Building Autonomous Agents, (Zhou et al., 2023)

Browsergym: a Gym Environment for Web Task Automation (Drouin et al., 2024)

# Recap: Architecture of an AI Agent





# What's special / difficult with Web Agents ?


## API Agents

- ❖ Observations: API call results, search history, user-uploaded images, chat history
- ❖ Actions: API calls, search calls, responses to the user
- ❖ Pros: Lower latency, lower risks
- ❖ Cons: needs appropriate APIs

## Web Agents

- ❖ Observations: what human would see + accessibility tree / raw DOM
- ❖ Actions: enter text in fields, clicks
- ❖ Pros: can do anything
- ❖ Cons: higher latency, higher risks

# Multi-modal Agents/ Multi-modal Agentic AI Applications




- ❖ Computer tasks often involve multiple apps and interfaces
- ❖ Powered by advancements in large vision-language-action models (VLA-Ms)
- ❖ Make digital interactions more accessible and vastly increase human productivity



# Web Agent Architecture

Model Architecture of an Interactive Agent:

- ❖ High-level Planning and Reasoning
- ❖ Observation Parsing
- ❖ Low-level Action Generation



# Some Challenges in Web Agents

MUST be able to deal with the **World WILD Web** !

- ❖ Long Context Understanding
  - ▶ HTML pages are complex, easily filling up > 100K tokens
- ❖ Long-term planning
  - ▶ Need to infer/ understand/ reason, set and meet complex objectives / constraints
  - ▶ Need to visit many websites, process man webpages to search for the right path ; errors can easily be cumulated and exponentially compounded
- ❖ Learning and adaptability
  - ▶ Messy HTML, Interactive/ Dynamic elements, New Features/ web programming construct/ languages!
- ❖ Strong Multimodality support/ understanding
  - ▶ Need to process and understand not only text feedback but also "pixel"s !
- ❖ Cost and efficiency
  - ▶ To be viable, Web Agents must produce more value than they cost !
- ❖ Safety and Alignment
  - ▶ Warn and protect you when your Web Agent logs into your Bank Account to wire your funds away !

# Web Agent Research Milestones


## Learning to Control Computers (DM)

- Control computers w/ keyboard & mouse from NL instructions
- MiniWoB++ through RL with computer-human interactions



## WebArena (CMU)

- Realistic benchmark, 812 tasks, 6 domains
- Long-horizon tasks
- Best GPT-4: 11% solve rate vs 78% for humans




## VisualWebArena

- Benchmark that needs visual comprehension
- Test visual & reasoning skills of web agents
- 910 tasks, 3 domains



## WebLINX (McGill)

- Conversational web agent navigation
- 2337 expert demos on 155 real-world websites
- Visual models not best; fine-tuning is key



2017

2021


2022

2023

2024

## World of Bits (WoB)

- First widely available web benchmark
- Simplified tasks
- 100 tasks
- Can be solved by RL



## WebGPT (OpenAI)

- Fine-tuned GPT-3 for QA with web browsing
- Evaluated on "Explain Like I'm 5" Reddit Qs + TruthfulQA dataset




## Mind2Web (Ohio)

- Benchmark of realistic web tasks from NL
- Interaction traces
- 2,350 tasks from 137 websites, 31 domains



## WebAgent (Google)

- Combine 2 LLMs to simplify huge HTML, plan solution, create code talking to web browser; **no pixels**
- MiniWoB & Mind2Web



## WebVoyager (Teng)

- Completes tasks on real websites using **textual+visual** inputs
- New benchmark: 15 websites, automatic GPT-4V-based eval.



# Web Agent Research Milestones (cont'd)

MAU)

mark,  
ains  
ks  
solve  
umans




## VisualWebArena

- Benchmark that needs **visual comprehension**
- Test visual & reasoning skills of web agents
- 910 tasks, 3 domains



## WebLINX (McGill)

- **Conversational** web agent navigation
- 2337 expert demos on 155 real-world websites
- Visual models not best; fine-tuning is key



## WorkArena (ServiceNow)

- Basic tasks that a knowledge worker must carry out
- Implemented on the ServiceNow platform



## OSWorld

- 369 computer tasks of real web and desktop apps in open domains
- OS file I/O + workflows spanning multiple applications




## WorkArena++ (ServiceNow)

- Compositional tasks with much higher difficulty than WorkArena
- Today's best models get single-digit performance, with huge room for improvement

2024

## WebAgent (Google)

- Combine 2 LLMs to simplify huge HTML, plan solution, create code talking to web browser; **no pixels**
- MiniWoB & Mind2Web



## WebVoyager (Tenor)

- Completes tasks on real websites using **textual+visual** inputs
- New benchmark: 15 websites, automatic GPT-4V-based eval.



## WebCanvas (CMU)

- Handles dynamic web
- Mind2Web-Live, a refined Mind2Web: 542 tasks, 2439 evaluation states



## AssistantBench

- Diverse web tasks: search, navigation, data extraction, interaction
- 214 tasks that can be auto-evaluated



## NNetNav (Stanford)

- Training web agents entirely through synthetic demos
- Web trajectory rollouts are processed by an LLM to be retroactively labeled into instruction




# Many Web Agent Benchmarks but lack unification

- ❖ MiniWoB++ (Shi et al., 2017; Liu et al., 2018) 125 tasks
- ❖ WebShop (Yao, Chen et al., 2022) 12087 tasks
- ❖ WebArena (Zhou et al., 2023) 812 tasks
- ❖ VisualWebArena (Koh et al., 2024) 910 tasks
- ❖ WebLINX (Lù et al., 2024) 2300 tasks
- ❖ WebCanvas (Pan et al., 2024) 438 tasks
- ❖ WebVoyager (He et al., 2024) 643 tasks
- ❖ AssistantBench (Yoran et al., 2024) 214 tasks
- ❖ WorkArena++ (ServiceNow Research, 2024) 682 tasks

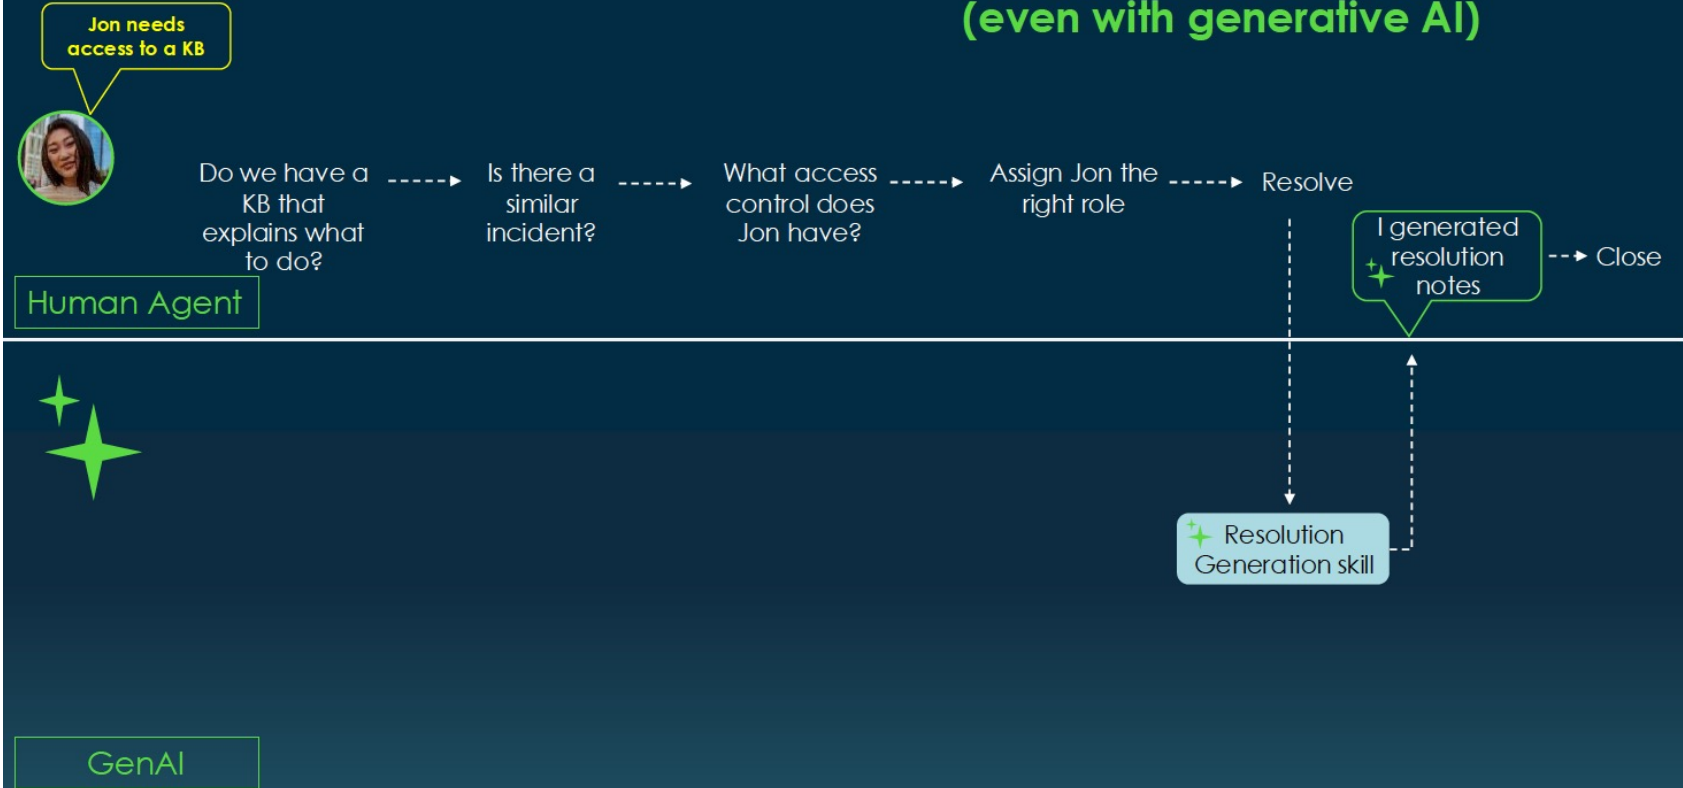


# Robotic Process Automation (RPA) w/ Physical Agents




# Impact of AI Agents on Job Nature

## Today's Enterprise Workflows Remain Quite Manual (even with generative AI)




# Impact of AI Agents on Job Nature (cont'd)

## Access issue – With AI Agents






# Automation in Enterprise Workflows



# Current Target Area of Web Agents



# Common Failure Modes of Current Web Agents and Future Directions

## ❖ Lack Long-Horizon Reasoning and Planning

- ▶ Models oscillate between two webpages, or get stuck in a loop
- ▶ Correctly performing tasks but undoing them
- ▶ Agents tend to stop exploration / execution too early

## How to Close the Gap ?

- ❖ Allow agent to Search, execute and coordinate multiple instances in parallel and ask for clarifications/confirmations
- ❖ Develop Stronger Multimodal, i.e. Vision-Language-Code Models
- ❖ Identify the appropriate level of abstraction for Agents (HTML/ Screenshots / APIs)



# GAIA: A Benchmark For General AI Assistants

<https://huggingface.co/spaces/gaia-benchmark/leaderboard>

Motivation: evaluating new AI systems requires to rethink benchmarks

- LLMs open the way to general purpose systems.
- Evaluating such systems is an open problem. For example, LLMs break AI benchmarks at an ever-increasing rate (Kiehl et al., 2023).
- Current trend: build more challenging benchmarks with tasks that are ever more difficult for humans. Example: MMLU (Hendrycks et al., 2021), close to be solved.
- More generally, open-ended generation difficult to evaluate:
  - Human evaluation: what about AI generated books? AI generated solution to complex and rare math problem?
  - Model evaluation: needs stronger model, subtle biases (Zheng et al., 2023).
- Alternatively, AI systems could be asked to solve conceptually simple tasks that require accurate execution of complex sequences of actions with large combinatorial spaces, with a unique answer easy to validate. Analogous to Proof of Work (Jakobsson and Juels, 1999).

## Our proposal: GAIA

- Benchmark for AI systems, **realistic general assistant questions**.
- 466 questions designed and annotated by humans.
- Text-based, sometimes come with a file such as images or spreadsheets.
- Questions require **fundamental abilities**: reasoning, multi-modality, code, tool use, various file type understanding.

**Level 1**

**Question:** What was the actual enrollment count of the clinical trial on H. pylori in acne vulgaris patients from Jan-May 2018 as listed on the NIH website?

**Ground truth:** 90

**Level 2**



**Question:** If this whole pint is made up of ice cream, how many percent above or below the US federal standards (or butterfat content) is it when using the standards as reported by Wikipedia in 2020? Answer as + or - a number rounded to one decimal place.

**Ground truth:** +4.6%

**Level 3**

**Question:** In NASA's Astronomy Picture of the Day on 2006 January 21, two astronauts are visible, with one appearing much smaller than the other. As of August 2023, out of the astronauts in the NASA Astronaut Group that the smaller astronaut was a member of, which one spent the least time in space, and how many minutes did he spend in space, rounded to the nearest minute? Exclude any astronauts who did not spend any time in space. Give the last name of the astronaut, separated from the number of minutes by a semicolon. Use commas as thousands separators in the number of minutes.

**Ground truth:** White; 5,876


Figure 1: Sample GAIA questions. Completing the tasks requires fundamental abilities such as reasoning, multi-modality handling, or tool use proficiency. Answers are unambiguous and by design unlikely to be found in plain text in training data. Some questions come with additional evidence, such as images, reflecting real use cases and allowing better control on the questions.



## Design choices

- Conceptually simple questions for humans, rooted in the real world, challenging for AI systems. Focus on fundamental abilities rather than specialized skills (Chollet, 2019).
- Interpretability:** small number of questions, conceptually simple so the results can be inspected.
- Robustness against contamination/memorization:** answers need to be extracted from the internet or documents; the map question-answer is not present in plain text online.
- Ease of use:** zero-shot prompting, single correct answer, allowing fast and factual evaluation.

## Composition of GAIA



- We rely on our annotators' metadata.
- Questions splitted into three levels of difficulty.
- Some questions come with additional files.

## Question creation process

- Question creation: annotators combine elements from one or more sources of truth (websites such as Wikipedia, or attached files), and answer. Effort to vary domains and question types.
- Question validation: two new annotators answer the newly created question. If both agree with original answer, question is validated. Otherwise, repaired or removed.

## Challenges associated with question creation

- Eliminating ambiguities from a question.
- Relying on the web: more ambiguities, webpage can change, robots.txt.

## AUTHORS

Grégoire Mialon (Meta GenAI),  
Clémentine Fourrier (Hugging Face),  
Thomas Wolf (Hugging Face),  
Yann LeCun (Meta FAIR),  
Thomas Scialom (Meta GenAI)

## LLMs results on GAIA

**System prompt:** You are a general AI assistant. I will ask you a question. You need to finish your answer with the following template: FINAL ANSWER: [YOUR FINAL ANSWER]. Your final answer should be a number, or as few words as possible.

**GAIA Question:** The attached Excel file contains the sales of menu items for a local fast-food chain. What were the total sales that the chain made from food (not including drinks)? Express your answer in USD, with a dollar sign and two decimal places.



**GP14 (Advanced Data Analysis mode):**

**Working:**

```
import pandas as pd

# Load the Excel file
file_path = "fast_data_uploaded.xlsx"
df = pd.read_excel(file_path)
```

**Working:**

```
# Calculate the total sales from food items (not including drinks i.e., Soda)
total_food_sales = df[["Burgers", "Hot Dogs", "Salads", "Fries", "Ice Cream"]
].sum().sum()

# Format the result to USD, with a dollar sign and two decimal places
total_food_sales_formatted = f"${total_food_sales:.2f}"
```

**FINAL ANSWER:** \$89,706.00

**Ground truth:** \$89,706.00 ✓




Figure 4: Left: scores per method and level. Right: scores per model and capabilities. As stated in the main text, GP14 + plugins score should be seen as an oracle since the plugins were chosen manually. Human score refers to the score obtained by our annotators when validating the questions.



## Discussion and Limitations

- Tools unlock many use cases. More work is needed and worth it.
- Reproducibility issues associated with closed models and plugins/tools.
- Static vs. dynamic benchmarks.
- Towards unified evaluation of generative models?
- Cost of designing unambiguous questions.

## References

- Douwe Kiela, Tristan Thrush, Kaveer Ethayarajh, and Amanpreet Singh. Plotting progress in AI. *Contextual AI Blog*, 2023. <https://contextual.ai/blog/plotting-progress>.
- Den Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Maseika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2022.
- Liammi Zheng, Wei Lu Chang, Ying Sheng, Siyuan Zhang, Zhengao Wu, Yinghan Zhang, Zhi Lin, Zhutian Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging llm-as-a-judge with mt-bench and chatbot arena, 2023.
- Markus Jakobsson and Jon Juels. *Proofs of Work and Bread Pudding Protocols (Extended Abstracts)*, pages 258–272. Springer US, Boston, MA, 1999. ISBN 978-0-387-35568-9. doi:10.1007/978-0-387-35568-9\_18. URL: [https://doi.org/10.1007/978-0-387-35568-9\\_18](https://doi.org/10.1007/978-0-387-35568-9_18).
- François Chollet. On the measure of intelligence, 2019.

# Recent Progress on Coding, Agents & Computer Use w/ Kimi



# Latest Progress on Coding, Agents & Computer Use



|                                                  | Opus 4.5         | Sonnet 4.5       | Opus 4.1         | Gemini 3 Pro     | GPT-5.1            |
|--------------------------------------------------|------------------|------------------|------------------|------------------|--------------------|
| Agentic coding<br>SWE-bench Verified             | 80.9%            | 77.2%            | 74.5%            | 76.2%            | 76.3%              |
|                                                  |                  |                  |                  |                  | 77.9%<br>Codex-Max |
| Agentic terminal<br>coding<br>Terminal-bench 2.0 | 59.3%            | 50.0%            | 46.5%            | 54.2%            | 47.6%              |
|                                                  |                  |                  |                  |                  | 58.1%<br>Codex-Max |
| Agentic tool use<br>t2-bench                     | Retail<br>88.9%  | Retail<br>86.2%  | Retail<br>86.8%  | Retail<br>85.3%  | —                  |
|                                                  | Telecom<br>98.2% | Telecom<br>98.0% | Telecom<br>71.5% | Telecom<br>98.0% | —                  |
| Scaled tool use<br>MCP Atlas                     | 62.3%            | 43.8%            | 40.9%            | —                | —                  |
| Computer use<br>OSWorld                          | 66.3%            | 61.4%            | 44.4%            | —                | —                  |
| Novel problem<br>solving<br>ARC-AGI-2 (Verified) | 37.6%            | 13.6%            | —                | 31.1%            | 17.6%              |
| Graduate-level<br>reasoning<br>GPQA Diamond      | 87.0%            | 83.4%            | 81.0%            | 91.9%            | 88.1%              |
| Visual reasoning<br>MMMU (validation)            | 80.7%            | 77.8%            | 77.1%            | —                | 85.4%              |
| Multilingual Q&A<br>MMMLU                        | 90.8%            | 89.1%            | 89.5%            | 91.8%            | 91.0%              |

<https://www.anthropic.com/news/claude-opus-4-5> ;

Date: Nov. 25, 2025